

CS 4410
Operating Systems

Synchronization
Spinlocks - Semaphores

Summer 2013
Cornell University

Today

- How can I synchronize the execution of multiple threads of the same process?
- Example
- Race condition
- Critical-Section Problem
- Spinlocks
- Semaphors
- Usage

Problem Context

- Multiple threads of the same process have:
 - Private registers and stack memory
 - Shared access to the remainder of the process “state”
- Preemptive CPU Scheduling:
 - The execution of a thread is interrupted unexpectedly.
- Multiple cores executing multiple threads of the same process.

Share Counting



- Mr Skroutz wants to count his \$1-bills.
- Initially, he uses one thread that increases a variable *bills_counter* for every \$1-bill.
- Then he thought to accelerate the counting by using two threads and keeping the variable *bills_counter* shared.

Share Counting



bills_counter = 0

- Thread A

```
while (machine_A_has_bills)
    bills_counter++
```

- Thread B

```
while (machine_B_has_bills)
    bills_counter++
```

***print** bills_counter*

- What it might go wrong?

Share Counting



- Thread A

$r1 = \text{bills_counter}$

$r1 = r1 + 1$

$\text{bills_counter} = r1$

- Thread B

$r2 = \text{bills_counter}$

$r2 = r2 + 1$

$\text{bills_counter} = r2$

- If *$\text{bills_counter} = 42$* , what are its possible values after the execution of one A/B loop ?

Shared counters

- One possible result: everything works!
- Another possible result: lost update!
- Called a “**race condition**”.

Race conditions

- Def: *a timing dependent error involving shared state*
 - It depends on how threads are scheduled.
- Hard to detect

Critical-Section Problem

bills_counter = 0

- Thread A

while (my_machine_has_bills)

- **enter critical section**

bills_counter++

- ***exit critical section***

- Thread B

while (my_machine_has_bills)

- **enter critical section**

bills_counter++

- ***exit critical section***

print bills_counter

Critical-Section Problem

- The solution should satisfy:
 - Mutual exclusion
 - Progress
 - Bounded waiting

- *enter section*
 - *critical section*
- *exit section*
 - *remainder section*

General Solution

- LOCK
- A process must acquire a lock to enter a critical section.
- Hardware or Software-based implementation



TestAndSet

- **Hardware** instruction.
- **Test** and **modify** the content of **one word atomically**.

```
boolean TestAndSet(boolean *target){  
    boolean rv = *target;  
    *target = TRUE;  
    return rv;  
}
```

Share Counting



bills_counter = 0
lock = FALSE

- Thread A

```
while (machine_A_has_bills){  
    while (TestAndSet(&lock))  
        ;  
    bills_counter++  
    lock = FALSE;  
}
```

- Thread B

```
while (machine_B_has_bills){  
    while (TestAndSet(&lock))  
        ;  
    bills_counter++  
    lock = FALSE;  
}
```

print bills_counter

```
boolean TestAndSet(boolean *target){  
    boolean rv = *target;  
    *target = TRUE;  
    return rv;  
}
```

Swap

- **Hardware** instruction.
- Swap the contents of two **words atomically**.

```
void Swap (boolean *a, boolean *b){  
    boolean temp = *a;  
    *a = *b;  
    *b = temp;  
}
```

Share Counting



bills_counter = 0
lock = FALSE

• Thread A

```
while (machine_A_has_bills){  
    keyA = TRUE;  
    while (keyA == TRUE)  
        Swap(&lock, &keyA);  
    bills_counter++;  
    lock = FALSE;  
}
```

• Thread B

```
while (machine_B_has_bills){  
    keyB = TRUE;  
    while (keyB == TRUE)  
        Swap(&lock, &keyB);  
    bills_counter++;  
    lock = FALSE;  
}
```

print bills_counter

```
void Swap (boolean *a, boolean *b){  
    boolean temp = *a;  
    *a = *b;  
    *b = temp;  
}
```

But...

- TestAndSet and Swap are complicated for application programmers to use.
- What is the alternative?

Semaphores

- Integer value
- Atomic operations
 - wait
 - signal

```
. . . . .  
: wait(S) {  
:     while S <= 0  
:         ;  
:     S--;  
: }  
. . . . .
```

```
. . . . .  
: signal(S) {  
:     S++;  
: }  
. . . . .
```

Share Counting



bills_counter = 0
S = 1

- Thread A

```
while (machine_A_has_bills){  
    wait(S);  
    bills_counter++  
    signal(S);  
}
```

- Thread B

```
while (machine_B_has_bills){  
    wait(S);  
    bills_counter++  
    signal(S);  
}
```

print bills_counter

```
wait(S) {  
    while S <= 0  
        ;  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

Spinlock

- This implementation of Semaphors, TestAndSet and Swap are spinlocks.
 - They require **busy waiting**.
- The waiting processes should **loop continuously** in the entry code.
- Valuable **CPU** cycles are **wasted**.
- Solution:
 - **Block** the waiting process.
 - **Signal** blocked process when the semaphore is “available”.

Semaphores

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore;
```

```
wait (semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```

```
signal (semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

Usage

- Binary semaphore (mutex)
 - Ranges between 0 and 1.
 - Ex. Only one process can access a resource.
- Counting semaphore
 - Ranges, usually, between 0 and N.
 - Ex. N resources and M processes that share the resources.
- Synchronization
 - Ranges between 0 and 1.
 - Ex. Process A should do task A_t after process B having done task B_t .

Today

- How can I synchronize the execution of multiple threads of the same process?
- Example
- Race condition
- Critical-Section Problem
- Spinlocks
- Semaphors
- Usage